



UPPSALA  
UNIVERSITET

# Contextualization and Containers

Salman Toor  
[salman.toor@it.uu.se](mailto:salman.toor@it.uu.se)



UPPSALA  
UNIVERSITET

# Images and snapshots

## Images

<input type="checkbox"/> ▼ <a href="#">Ubuntu 20.04 - 2020.12.01</a>		Image	Active	Public	No	RAW	2.20 GB	Launch ▼
<b>Name</b> Ubuntu 20.04 - 2020.12.01	<b>Visibility</b> Public							
<b>ID</b> ffecce9c-4c76-4fdc-ad76-ad1cc614f34c	<b>Protected</b> No							

## Snapshots

Snapshots		Displaying 99 items   <a href="#">Next »</a>						
	Groups	<input type="checkbox"/> Name	Description	Size	Status	Group Snapshot	Volume Name	Actions
	Group Snapshots	<input type="checkbox"/> <a href="#">snapshot for Group10_project_snapshot</a>	-	20GiB	Available	-	<a href="#">e054e9e0-03e6-4c17-9e4c-b3c3121041ee</a>	Create Volume ▼
	Network >	<input type="checkbox"/> <a href="#">snapshot for Group10_project_snapshot</a>	-	200GiB	Available	-	<a href="#">Group_10_volume</a>	Create Volume ▼
	Orchestration >	<input type="checkbox"/> <a href="#">MengfeiLiang_A3_ss</a>	-	20GiB	Available	-	<a href="#">MengfeiLiang_A3_vol</a>	Create Volume ▼
	Identity >	<input type="checkbox"/> <a href="#">snapshot for stefan_aslan_A3</a>	-	20GiB	Available	-	<a href="#">StefanAslanA3</a>	Create Volume ▼
		<input type="checkbox"/> <a href="#">snapshot for brent-hefty-A2-20</a>	-	20GiB	Available	-	<a href="#">brent-hefty-A2-volume</a>	Create Volume ▼
		<input type="checkbox"/> <a href="#">alma_backman_a1_snap</a>	-	20GiB	Available	-	<a href="#">alma_backman_v_a1</a>	Create Volume ▼
		<input type="checkbox"/> <a href="#">OK_TO_DELETE</a>	-	20GiB	Available	-	<a href="#">ef3b0dd1-5d18-4a27-9ef7-3a84b197379a</a>	Create Volume ▼



# Images and formats

- Cloud images are customized disks of OSes for private or public clouds
- Different formats are available:
  - raw: An unstructured disk image format (big in size)
  - vhd: VMware, Xen, Microsoft, VirtualBox, and others
  - vdi: Supported by VirtualBox, QEMU emulator.
  - iso: Archive format for the data contents of an optical disc
  - qcow2: Supported by the QEMU emulator that can expand dynamically and supports Copy on Write.
  - ...



# Contextualization

In cloud computing, contextualization means providing a customized computing environment

Or

Allows a virtual machine instance to learn about its cloud environment and user requirement (the 'context') and configure itself to run correctly



# Contextualization

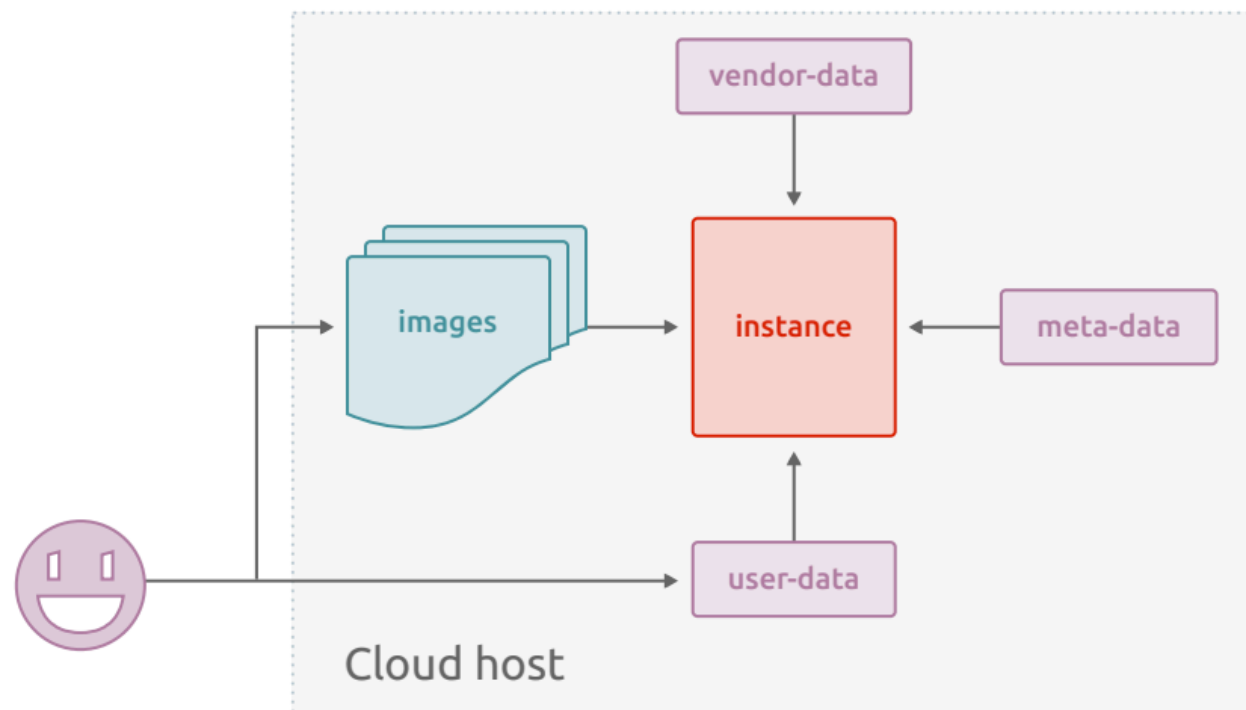
- Provide scalable solution
- No need to manage fat images
- Dynamic configuration
- Typically work in two layers
  - Meta-data : System information handled at cloud level
  - User-defined-data: User specific requirements/settings





# CloudInit

Cloud-init is a standard multi-distribution method for cross-platform cloud instance initialization. It is supported across all major public cloud providers, provisioning systems for private cloud infrastructure, and bare-metal installations.





UPPSALA  
UNIVERSITET

# CloudInit

```
#cloud-config
apt_update: true
apt_upgrade: true
packages:
  - cowsay
  - python-pip
  - python-dev
  - build-essential
  - cowsay
byobu_default: system

runcmd:
  - echo "export PATH=$PATH:/usr/games" >> /home/ubuntu/.bashrc
  - source /home/ubuntu/.bashrc
  - pip install Flask
  - python /home/ubuntu/cowsay-app.py &
```



UPPSALA  
UNIVERSITET

# Beyond Virtual Machines





# Containers

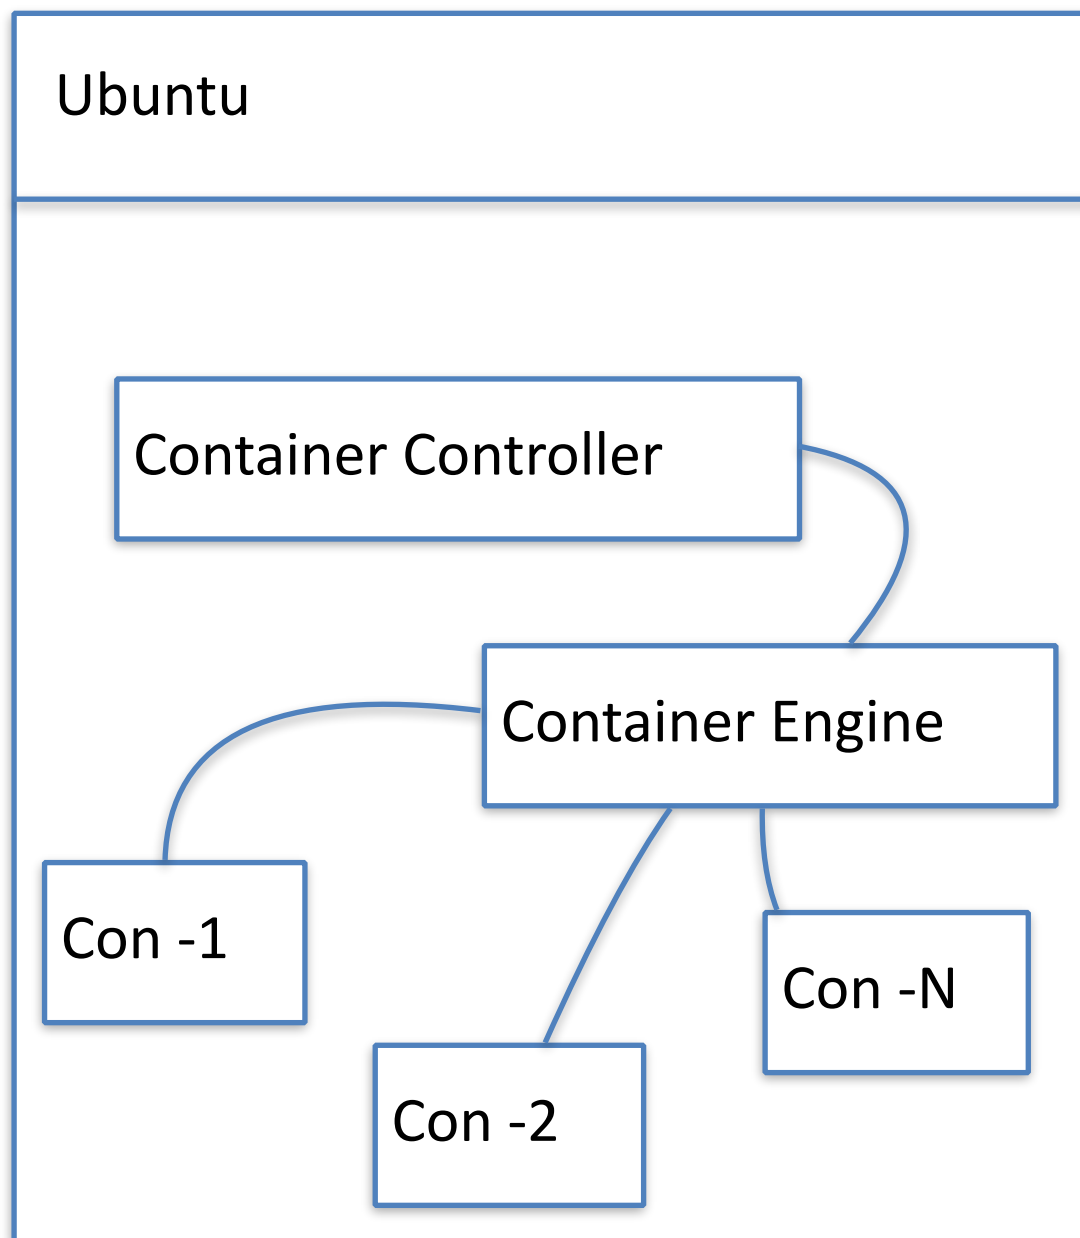
- OS level virtualization environment
  - Kernelspace is shared
  - Userspace is separate for each linux system (container)
- A lightweight alternative to Virtual Machines (VM)
- Shared same resources as host OS
- A simple model for packaging applications in Linux.



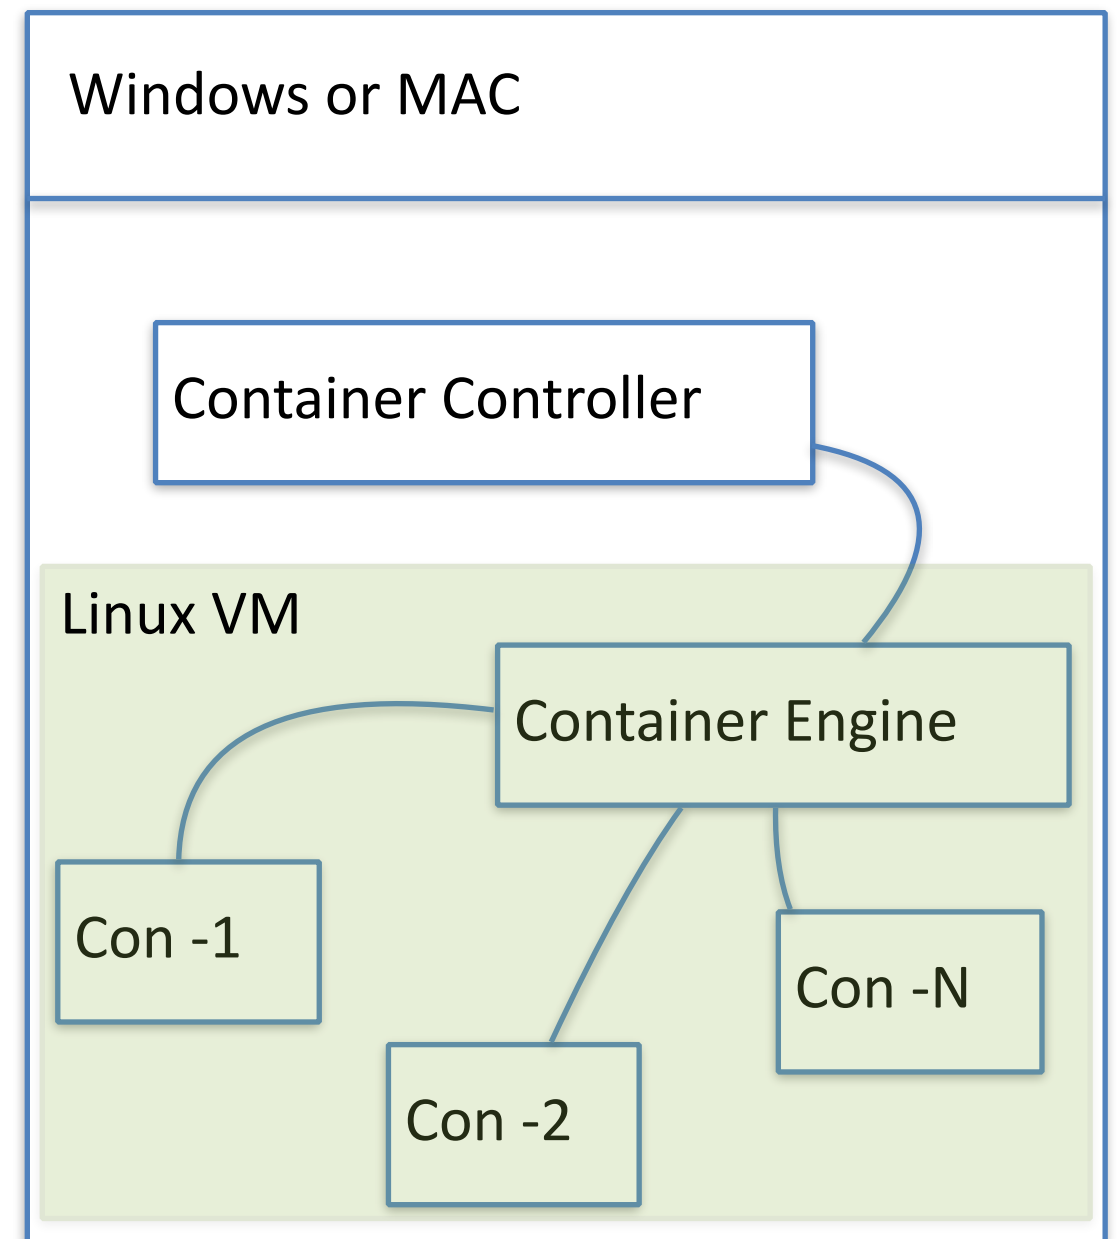
UPPSALA  
UNIVERSITET

# Basic illustration of containers

## Linux based host OS



## Non Linux host OS





# VM and containers

**It is important to understand that VMs and Containers should not be viewed as competitors**

## Virtual Machine

- Complete isolation
- Big in size
- High overhead
- Flexible support of multiple OSes
- Greater stability both for hypervisors and VMs
- Better security

## Containers

- Application level abstraction
- Lightweight
- Works well with Linux, limited support for Windows
- Weak security
- Significant management Overhead
- Not well suited for large applications
- Important for micro-services design



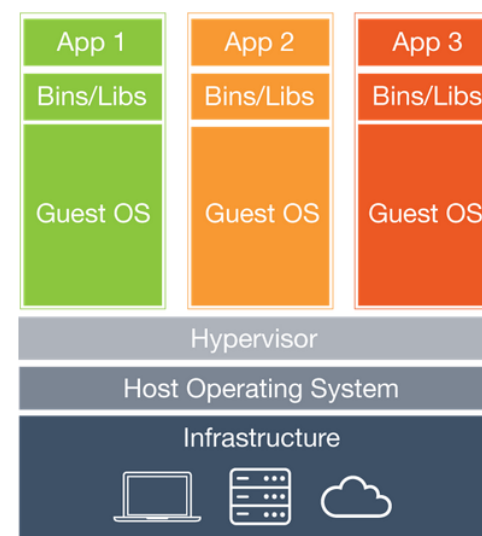
# LXC

- Linux Container (LXC) is an open source software
- Virtual environment based on separate memory, CPU, network, io etc
- Similar to the concept of *chroot*
- Used in most of the container based orchestration tools
- LXD is a newer version of LXC, advanced and stronger support for cloud plugins

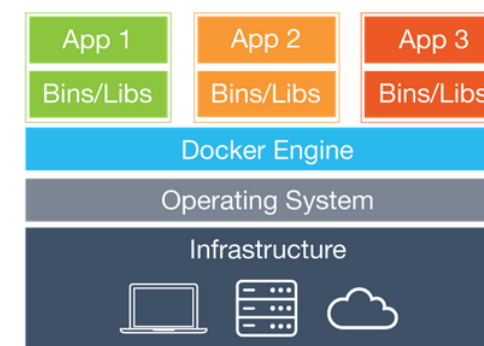


# Docker

- Docker package an application together with all its dependencies in the container
- Guarantees that it will always run the same regardless of the environment
- Container based orchestration tool
- Docker Hub, container registry
- Open source



Virtual Machines




Containers





# Docker framework



```
# docker ps
# docker images
# docker build
# docker run
# docker exec
# docker pull
# docker push
... •
```



# Dockerfile

- A `Dockerfile` is a text file that contains all commands, in order, needed to build a given image
- It adheres to a specific format and set of instructions
- Using `docker build`, users can create an automated build that executes several command-line instructions in succession

```
# our base image
FROM alpine:3.5

# Install python and pip
RUN apk add --update py2-pip

# upgrade pip
RUN pip install --upgrade pip

# install Python modules needed by the Python app
COPY requirements.txt /usr/src/app/
RUN pip install --no-cache-dir -r /usr/src/app/requirements.txt

# copy files required for the app to run
COPY app.py /usr/src/app/
COPY templates/index.html /usr/src/app/templates/

# tell the port number the container should expose
EXPOSE 5000

# run the application

CMD ["python", "/usr/src/app/app.py"]
```



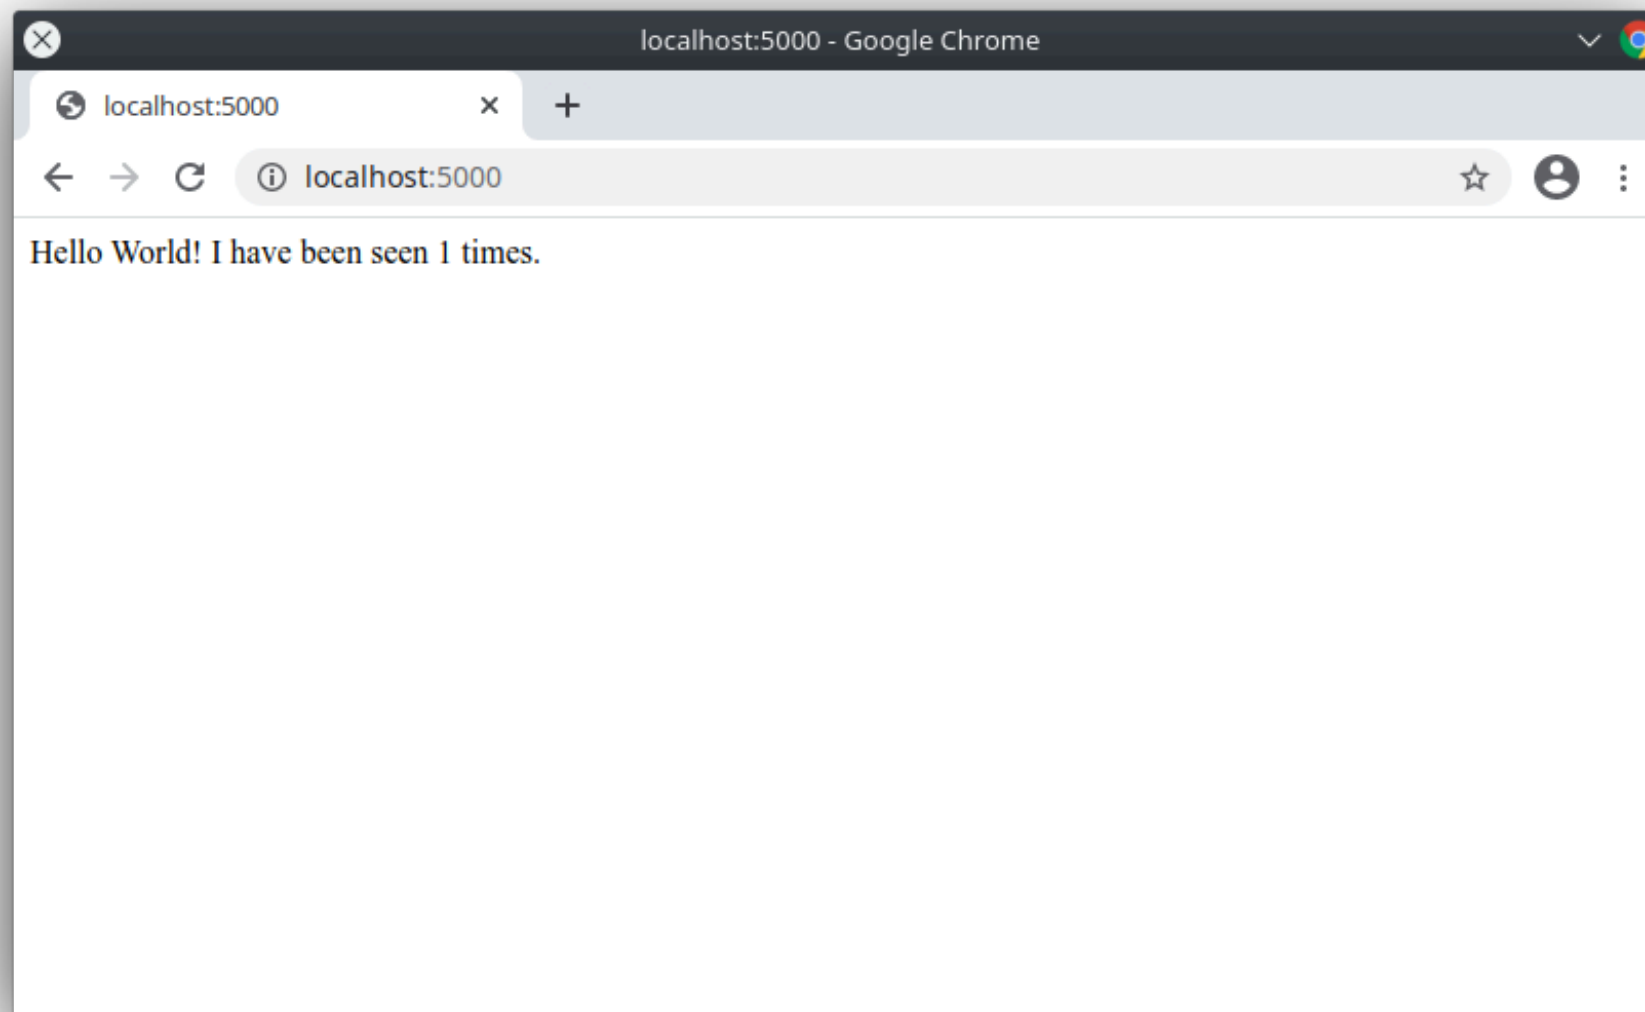
# Docker-compose

- `docker-compose` can be viewed as an automated multi-container workflow
- It belongs to the Docker family, used to define and run multi-container applications
- With *compose*, we use a YAML file to configure application's services and create all the app's services from that configuration



UPPSALA  
UNIVERSITET

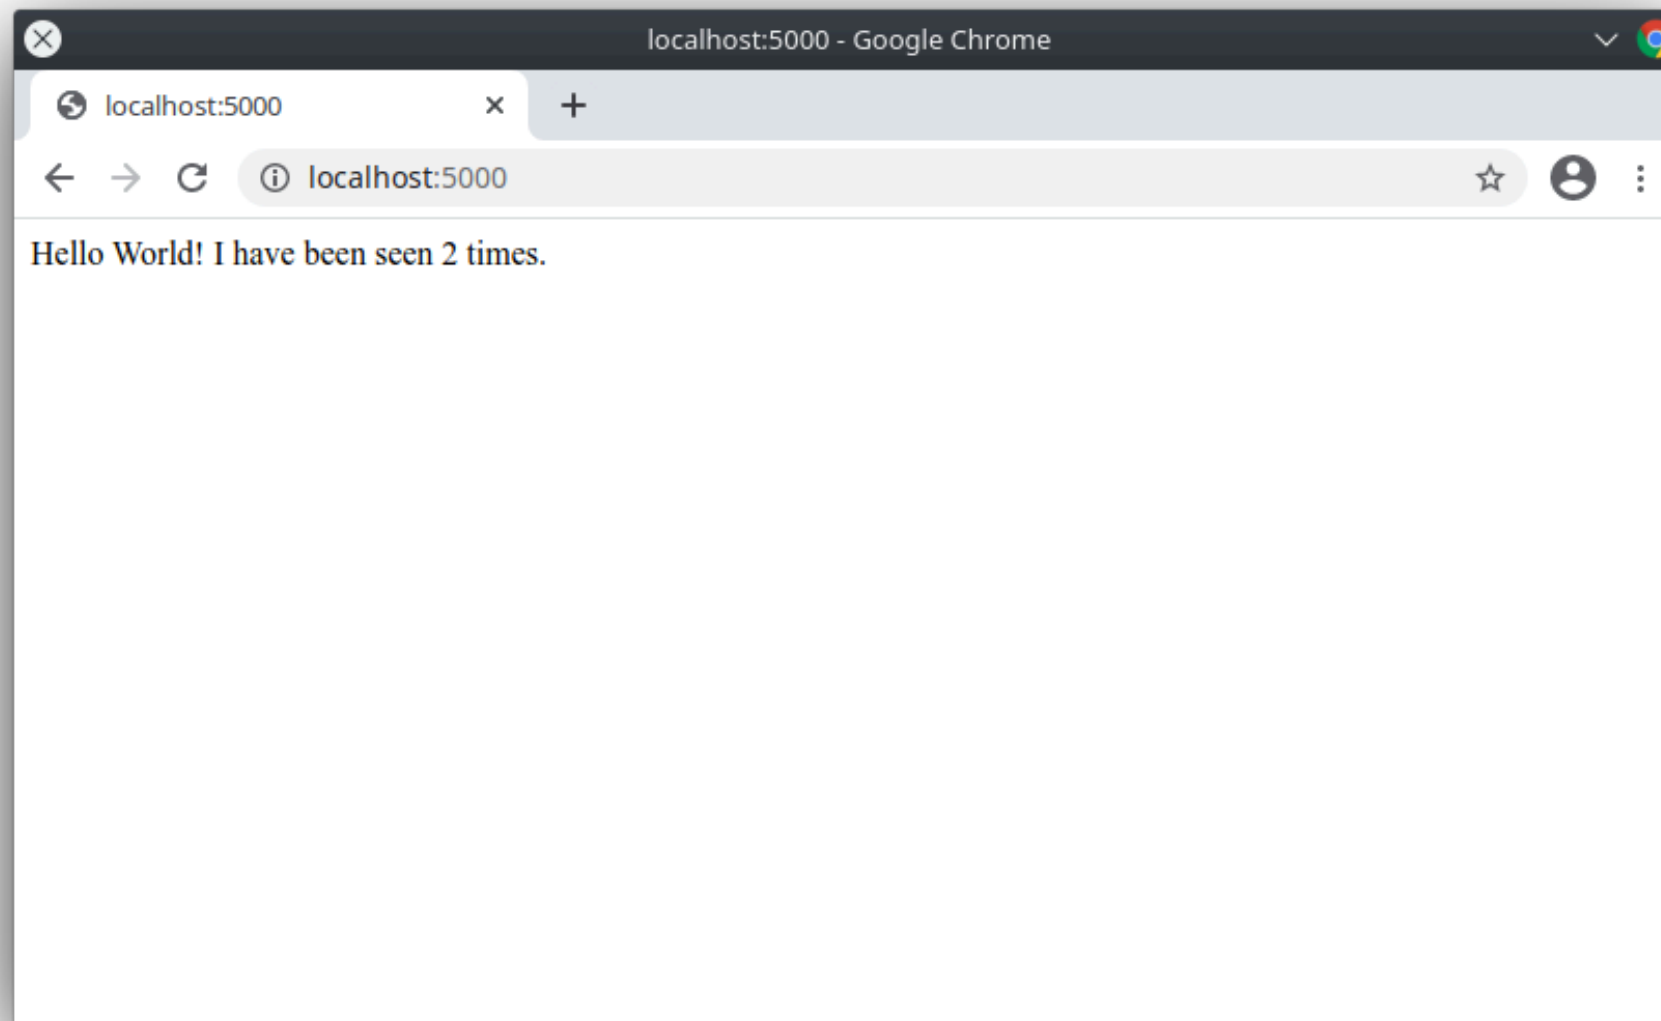
# Docker-compose example





UPPSALA  
UNIVERSITET

# Docker-compose example







# Docker-compose example

## app.py

```
import time

import redis
from flask import Flask

app = Flask(__name__)
cache = redis.Redis(host='redis', port=6379)

def get_hit_count():
    retries = 5
    while True:
        try:
            return cache.incr('hits')
        except redis.exceptions.ConnectionError as exc:
            if retries == 0:
                raise exc
            retries -= 1
            time.sleep(0.5)

@app.route('/')
def hello():
    count = get_hit_count()
    return 'Hello World! I have been seen {} times.\n'.format(count)
```



UPPSALA  
UNIVERSITET

# Docker-compose example

**requirements.txt**

```
flask  
redis
```



# Docker-compose example

## Dockerfile

```
FROM python:3.7-alpine
WORKDIR /code
ENV FLASK_APP=app.py
ENV FLASK_RUN_HOST=0.0.0.0
RUN apk add --no-cache gcc musl-dev linux-headers
COPY requirements.txt requirements.txt
RUN pip install -r requirements.txt
EXPOSE 5000
COPY . .
CMD ["flask", "run"]
```



UPPSALA  
UNIVERSITET

# Docker-compose example

## **docker-compose.yml**

```
version: "3.9"
services:
  web:
    build: .
    ports:
      - "5000:5000"
  redis:
    image: "redis:alpine"
```



UPPSALA  
UNIVERSITET

# Docker-compose example

**command line interface**

```
# docker-compose up .
```





# Docker-compose example

## requirements.txt

```
flask  
redis
```

3

## Dockerfile

```
FROM python:3.7-alpine  
WORKDIR /code  
ENV FLASK_APP=app.py  
ENV FLASK_RUN_HOST=0.0.0.0  
RUN apk add --no-cache gcc musl-dev linux-headers  
COPY requirements.txt requirements.txt  
RUN pip install -r requirements.txt  
EXPOSE 5000  
COPY . .  
CMD ["flask", "run"]
```

2

## docker-compose.yml

```
version: "3.9"  
services:  
  web:  
    build: .  
    ports:  
      - "5000:5000"  
  redis:  
    image: "redis:alpine"
```

## command line interface

```
# docker-compose up .
```

1

## app.py

```
import time  
  
import redis  
from flask import Flask
```



# Docker Storage

- Containers are ephemeral (data is lost when the container stops).
- Solutions for persistent data:
  - **Volumes:** Managed by Docker (stored in /var/lib/docker/volumes).
  - **Bind Mounts:** Link container paths to host machine directories.
  - **tmpfs Mounts:** Store data in host memory (non-persistent).

```
# Create a volume
docker volume create my_volume

# Mount a volume to a container
docker run -d -v my_volume:/app/data nginx

# Bind mount (host dir → container)
docker run -d -v /host/path:/container/path nginx
```



# Docker Storage

- Why Use Volumes?
  - Decouple storage from containers.
  - Back up/restore easily.
  - Share data between containers.

```
# Container 1 writes to a shared volume
docker run -d -v shared_vol:/data --name writer alpine sh -c "echo 'Hello' > /data/file.txt"

# Container 2 reads from the same volume
docker run -d -v shared_vol:/data --name reader alpine cat /data/file.txt
```



# Docker Networking Basics

- **Default Networks:**
  - **Bridge:** Default network for containers (isolated communication).
  - **Host:** Bypasses Docker's network isolation (container uses host's network).
  - **None:** No networking (only loopback).

```
# Create a network
docker network create my_network

# Run containers on the same network
docker run -d --name web --network my_network nginx
docker run -d --name app --network my_network alpine ping web
```



# Networking Use Case

- Multi-Service App with Isolated Networks







# Networking Use Case

- Multi-Service App with Isolated Networks

- Create a network.

```
# Create a frontend network (public-facing)
docker network create frontend --driver bridge

# Create a backend network (private, for DBs)
docker network create backend --driver bridge
```

- Services with Network Isolation

```
# PostgreSQL (only on backend network)
docker run -d \
  --name postgres \
  --network backend \
  -e POSTGRES_PASSWORD=secret \
  -v pg_data:/var/lib/postgresql/data \
  postgres:13

# Redis (only on backend network)
docker run -d \
  --name redis \
  --network backend \
  redis:6

# Flask App (connected to both networks)
docker run -d \
  --name flask-app \
  --network frontend \
  --network backend \ # Connects to both!
  -p 5000:5000 \
  -e DB_HOST=postgres \
  -e REDIS_HOST=redis \
  my-flask-app-image
```



# Networking Use Case

- Multi-Service App with Isolated Networks
  - Verify Connectivity

```
# Test app → PostgreSQL
docker exec -it flask-app ping postgres # Should work

# Test app → Redis
docker exec -it flask-app ping redis    # Should work

# Test external access (only Flask exposed)
curl http://localhost:5000
```



# Docker Secrets

- Securing Sensitive Data in Docker
  - Problem:
    - Hardcoding secrets (passwords, API keys) in images or env files is unsafe.
    - Containers can leak secrets via logs or inspect tools.
  - Solutions:
    - Docker Secrets (Native for Swarm, but usable in Compose).
    - Environment Variables (Limited security).
    - Secret Mounts (Files injected at runtime).
    - Third-party Tools (HashiCorp Vault, AWS Secrets Manager).



# Docker Secrets

- Secrets in Docker
  - Docker Compose Example
  - Start PostgreSQL with Volume.
- Connect an app container.

```
services:
  db:
    image: postgres
    secrets:
      - db_password
    environment:
      POSTGRES_PASSWORD_FILE: /run/secrets/db_password

secrets:
  db_password:
    file: ./secrets/db_password.txt # External file (not in repo)
```

```
# Create a secret (Swarm mode only)
echo "mysecret123" | docker secret create db_password -
```

```
# Use a bind mount (temporary workaround)
docker run -v $(pwd)/secrets:/secrets alpine cat /secrets/db_password
```





# Docker Secrets

- Secrets in Docker
  - **Never** hardcode secrets in **Dockerfiles** or source code.
  - Prefer **Docker Secrets** or external tools (Vault) for production.
  - **Audit** secret usage (e.g., `docker secret ls`).





# Container softwares

- OpenVZ
- Virtuozzo (Linux and Windows)
- Solaris Containers (Solaris)
- Spoon (Windows)
- VMware ThinApp (Windows)



# Docker Swarm

Docker swarm is a container orchestration tool that allows users to manage multiple containers deployed across multiple host machines.

